

# Artificial Intelligence - Week 1 Homework

## Question 1

We start by including the Gaussian distribution used to draw the  $X$

$$\begin{aligned} p(x, y) &= \sum_{i \in (1,2)} p(x, y, c(x) = i) \\ &= \sum_{i \in (1,2)} p(x, y \mid c(x) = i) \cdot p(c(x) = i) \end{aligned}$$

as  $y$  and  $x$  are independent, but both are dependent on  $c(x)$ , it can be written as follows.

$$p(x, y) = \sum_{i \in (1,2)} p(x \mid c(x) = i) \cdot p(y \mid c(x) = i) \cdot p(c(x) = i)$$

where

$$p(x \mid c(x)) = \begin{cases} \sim N\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) & \text{if } c(x) = 1 \\ \sim N\left(\begin{bmatrix} 2.5 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) & \text{if } c(x) = 2 \end{cases}$$

$$p(y = 0 \mid c(x)) = \begin{cases} 0.8 & \text{if } c(x) = 1 \\ 0.7 & \text{if } c(x) = 2 \end{cases}$$

$$\text{and } p(c(x) = 1) = 0.8$$

Thus the distribution becomes

$$\begin{aligned} p(x, y) &= p(x \mid c(x) = 1) \cdot p(y \mid c(x) = 1) \cdot p(c(x) = 1) + p(x \mid c(x) = 2) \cdot p(y \mid c(x) = 2) \cdot p(c(x) = 2) \\ &= \begin{cases} p(x \mid c(x) = 1) \cdot 0.8 \cdot 0.8 + p(x \mid c(x) = 2) \cdot 0.7 \cdot 0.2 & \text{if } y = 0 \\ p(x \mid c(x) = 1) \cdot 0.2 \cdot 0.8 + p(x \mid c(x) = 2) \cdot 0.3 \cdot 0.2 & \text{if } y = 1 \end{cases} \\ &= \begin{cases} 0.64 \cdot X \sim N\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) + 0.14 \cdot X \sim N\left(\begin{bmatrix} 2.5 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) & \text{if } y = 0 \\ 0.16 \cdot X \sim N\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) + 0.06 \cdot X \sim N\left(\begin{bmatrix} 2.5 \\ 0 \end{bmatrix}, \begin{bmatrix} 9 & 0 \\ 0 & 1 \end{bmatrix}\right) & \text{if } y = 1 \end{cases} \\ p(x, y) &= \begin{cases} X \sim N\left(\begin{bmatrix} 0.35 \\ 0 \end{bmatrix}, \begin{bmatrix} 3.8628 & 0 \\ 0 & 0.4292 \end{bmatrix}\right) & \text{if } y = 0 \\ X \sim N\left(\begin{bmatrix} 0.15 \\ 0 \end{bmatrix}, \begin{bmatrix} 8.29 \times 10^{-4} & 0 \\ 0 & 9.21 \times 10^{-5} \end{bmatrix}\right) & \text{if } y = 1 \end{cases} \end{aligned}$$

## Question 2

**1. Show that the solution for optimum weight  $\tilde{w}$  still takes the similar form.**

Starting from the loss function,

$$\begin{aligned}
 Loss &= \frac{1}{2N} \sum_{i=1}^n (y_i - \phi(x_i) \cdot w)^2 \\
 Loss &= \frac{1}{2N} (Y - \Phi \cdot w)^T (Y - \Phi \cdot w) \quad (\text{converting to matrix}) \\
 \frac{dLoss}{dw} &= \frac{1}{2N} (-2\Phi^T) (Y - \Phi \cdot w) \quad (\text{differentiation}) \\
 0 &= \Phi^T \cdot Y - \Phi^T \cdot \Phi \cdot w \quad (\text{minimizing}) \\
 \Phi^T \cdot \Phi \cdot w &= \Phi^T \cdot Y \\
 w &= (\Phi^T \cdot \Phi)^{-1} (\Phi^T \cdot Y)
 \end{aligned}$$

**2. Write  $f(z)$  in terms of  $K$  and  $Y$ , where  $k(x, z) = \phi(x) \cdot \phi(z)$**

using  $w = \Phi^T \cdot v$ ,

$$\begin{aligned}
 Loss &= \frac{1}{2N} (Y - \Phi \cdot \Phi^T \cdot v)^T (Y - \Phi \cdot \Phi^T \cdot v) \\
 \frac{dLoss}{dv} &= \frac{1}{2N} (-2(\Phi \cdot \Phi^T)^T) (Y - \Phi \cdot \Phi^T \cdot v) \\
 0 &= (\Phi \cdot \Phi^T)^T (Y - \Phi \cdot \Phi^T \cdot v)
 \end{aligned}$$

since  $(\Phi \cdot \Phi^T)^T = (\Phi^T)^T \cdot (\Phi)^T = \Phi \cdot \Phi^T$ ,

$$\begin{aligned}
 0 &= \Phi \cdot \Phi^T (Y - \Phi \cdot \Phi^T \cdot v) \\
 0 &= \Phi \cdot \Phi^T \cdot Y - \Phi \cdot \Phi^T \cdot \Phi \cdot \Phi^T \cdot v \\
 \Phi \cdot \Phi^T \cdot \Phi \cdot \Phi^T \cdot v &= \Phi \cdot \Phi^T \cdot Y \\
 v &= (\Phi \cdot \Phi^T \cdot \Phi \cdot \Phi^T)^{-1} \Phi^T \cdot \Phi \cdot Y
 \end{aligned}$$

where

$$\begin{aligned}
 \Phi \cdot \Phi^T &= \begin{bmatrix} \phi(x_1) \\ \phi(x_2) \\ \vdots \\ \phi(x_n) \end{bmatrix} [\phi(x_1) \quad \phi(x_2) \quad \dots \quad \phi(x_n)] \\
 &= \begin{bmatrix} \phi(x_1) \cdot \phi(x_1) & \phi(x_1) \cdot \phi(x_2) & \dots & \phi(x_1) \cdot \phi(x_n) \\ \phi(x_2) \cdot \phi(x_1) & \phi(x_2) \cdot \phi(x_2) & \dots & \phi(x_2) \cdot \phi(x_n) \\ \vdots & \vdots & \ddots & \vdots \\ \phi(x_n) \cdot \phi(x_1) & \phi(x_n) \cdot \phi(x_2) & \dots & \phi(x_n) \cdot \phi(x_n) \end{bmatrix} \\
 &= \begin{bmatrix} \kappa(x_1, x_1) & \kappa(x_1, x_2) & \dots & \kappa(x_1, x_n) \\ \kappa(x_2, x_1) & \kappa(x_2, x_2) & \dots & \kappa(x_2, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ \kappa(x_n, x_1) & \kappa(x_n, x_2) & \dots & \kappa(x_n, x_n) \end{bmatrix} = K \\
 \Phi \cdot \Phi^T &= K
 \end{aligned}$$

therefore,

$$\begin{aligned}
 v &= (K \cdot K)^{-1} (K \cdot Y) \\
 f(z) &= \phi(z) \cdot \Phi^T \cdot v \\
 f(z) &= [\phi(z) \cdot \phi(x_1) \quad \phi(z) \cdot \phi(x_2) \quad \dots \quad \phi(z) \cdot \phi(x_n)] \cdot v \\
 f(z) &= [\kappa(z, x_1) \quad \kappa(z, x_2) \quad \dots \quad \kappa(z, x_n)] \cdot v \\
 f(z) &= [\kappa(z, x_1) \quad \kappa(z, x_2) \quad \dots \quad \kappa(z, x_n)] \cdot (K \cdot K)^{-1} (K \cdot Y)
 \end{aligned}$$

## Question 3

### I. Linear Features

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
```

a) Write a routine that loads a data file and returns a matrix  $X$  containing all  $x_i$  as rows, and a vector  $y$  containing all  $y_i$

```
In [2]: def load_dataset(filename):
    x = []
    y = []
    with open(filename, 'r') as inputfile:
        for inline in inputfile:
            data = inline.strip('\n').split()
            # filling y
            y.append(float(data[-1]))
            # filling x
            x.append([float(xi) for xi in data[:-1]])
    X = np.asarray(x)
    Y = np.asarray(y)

    return X, Y

# testing
x1d, y1d = load_dataset('dataLinReg1D.txt')
x2d, y2d = load_dataset('dataLinReg2D.txt')
# print(x1d)
```

b) Write a routine that takes the raw  $X$  as input and returns a new  $X$  with a '1' appended to each row. This routine simply computes the linear features including the constant '1'. This routine can later be replaced by others to work with non-linear features.

*Note: this notebook will use appended 1 (instead of pre-pend 1)*

```
In [3]: def linear_transform(X, D):
    return np.append(X, np.ones((X.shape[0],1)), axis=1)

# testing
x1d_appended = linear_transform(x1d, x1d)
# print(x1d_appended)
```

c) Write a routine that returns the optimal  $w$  from  $X$  and  $y$  - analytically, not by gradient descent.

```
In [4]: def ridge_regression(X, Y, lamb=1e-3):
    to_inv = np.matmul(X.T, X) + lamb * np.eye(X.shape[1])
    # print("to_inv:", to_inv.shape, np.linalg.det(to_inv))
    pre = np.linalg.inv(to_inv)
    post = np.matmul(X.T, Y)
    return np.matmul(pre, post)

# testing
w = ridge_regression(x1d_appended, y1d)
print(w)

[ 0.54359668 -0.79145962]
```

d) Generate some test data points (along a grid) and collect them in a matrix  $Z$ . Apply routine b) to compute features. Compute the predictions  $y = Zw$  (simple matrix multiplication) on the test data and plot it.

```


In [5]: def plot_z_matrix(X, Y, count=50):
        Xi = linear_transform(X, X)
        w = ridge_regression(Xi, Y)
        Z = np.random.randn(count, 1)
        Zi = linear_transform(Z, Z)
        Yz = np.matmul(Zi, w)

        plt.plot(X, Y, '.r')
        plt.plot(Z, Yz, '.b')

        # testing on 1d
        print("Note: Blue points are from the samples Z, Red points are from the dataset")
        plot_z_matrix(x1d, y1d)

```

Note: Blue points are from the samples Z, Red points are from the dataset



## II. Cross-validation

```

In [6]: x2d, y2d = load_dataset('dataLinReg2D.txt')

def rbf(x, y, alp=1e-3):
    norm = np.linalg.norm((x-y))
    return np.exp(-(norm/alp)**2)

def rbf_transform(X, D, alp=1e-3):
    rbfX = np.zeros((X.shape[0], D.shape[0]+1))

    for i in range(X.shape[0]):
        for j in range(D.shape[0]):
            rbfX[i, j] = rbf(X[i,:], D[j,:], alp=alp)
        rbfX[i, D.shape[0]] = 1

    return rbfX

# testing
rx1d = rbf_transform(x1d, x1d)

```

```

In [7]: def loss(X, Y, w, lamb):
#       Yi = np.matmul(X, w)
       Yi = np.matmul(X,w)
       lse = 0.5 * np.linalg.norm(Yi - Y)**2
       ridge = lamb * np.linalg.norm(w)**2
       return lse + ridge

def cross_validate(X, Y, transform, shuffled_indices, folds=5, lamb=1e-3):
    N = X.shape[0]
    pn = int(np.ceil(N / folds))

    idx_pieces = [shuffled_indices[i*pn:min((i+1)*pn,N)] for i in range(folds)]

    Ws = []
    crossval_losses = []

    for i in range(folds):
#         print("crossval",i)
        idx_p = idx_pieces[:]
        del idx_p[i]

        Xt = X[np.concatenate(idx_p)]
        Xtrain = transform(Xt, Xt)
        Ytrain = Y[np.concatenate(idx_p)]

        Xv = X[idx_pieces[i]]
        Xval = transform(Xv, Xt)
        Yval = Y[idx_pieces[i]]

        w = ridge_regression(Xtrain, Ytrain, lamb=lamb)
#         print("w",w)
        Ws.append(w)

        cross_loss = loss(Xval, Yval, w, lamb)
        crossval_losses.append(cross_loss)

    return np.mean(crossval_losses)

# testing

x1d, y1d = load_dataset('dataLinReg1D.txt')

# randomize, then split
shuffled_indices = np.arange(x1d.shape[0])
np.random.shuffle(shuffled_indices)

CVLoss = cross_validate(x2d, y2d, linear_transform, shuffled_indices)
print('Linear:',CVLoss)
# ERROR - handle rbf pls
CVLoss = cross_validate(x2d, y2d, rbf_transform, shuffled_indices)
print('RBF:',CVLoss)

```

```

Linear: 0.11468528441796835
RBF: 17.336489677609876

```

```
In [8]: # lamb_list = [1e-10, 1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1]

lamb_list = [1e-4, 1e-3, 1e-2, 1e-1, 1, 10]

opt_lamb = []
opt_cval = []

ropt_lamb = []
ropt_cval = []

for k in range(10):
    cval_list = []
    rcval_list = []

    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)

    for l in lamb_list:
        # print(">> lambda", l)
        cval_list.append(cross_validate(x2d, y2d, linear_transform, shuffled_indices, lamb=l))
        rcval_list.append(cross_validate(x2d, y2d, rbf_transform, shuffled_indices, lamb=l))

    opt_idx = cval_list.index(min(cval_list))
    opt_lamb.append(lamb_list[opt_idx])
    opt_cval.append(min(cval_list))

    ropt_idx = rcval_list.index(min(rcval_list))
    ropt_lamb.append(lamb_list[ropt_idx])
    ropt_cval.append(min(rcval_list))
```

```
In [9]: def bin_opt_lamb(opt, llist=lamb_list):
    opt_lamb_bin = []
    for l in llist:
        bin_count = 0
        for i in opt:
            if i == l:
                bin_count += 1
        opt_lamb_bin.append(bin_count)
    return opt_lamb_bin

op_lbin = bin_opt_lamb(opt_lamb)
rop_lbin = bin_opt_lamb(ropt_lamb)


def plot_bin(lbin, title, llist=lamb_list):
    # print(lbin)
    plt.figure(figsize=(10, 3))
    plt.bar([i for i in range(len(llist))], lbin, align='center', alpha=1)
    plt.title(title)
    plt.xticks([i for i in range(len(llist))], llist)
    plt.show()

print("Linear")
print('optimal lambda:', opt_lamb)
# print('cval losses:', opt_cval)
plot_bin(op_lbin, "Linear")

print("RBF")
print('optimal lambda:', ropt_lamb)
# print('cval losses:', ropt_cval)
plot_bin(rop_lbin, "RBF")
```

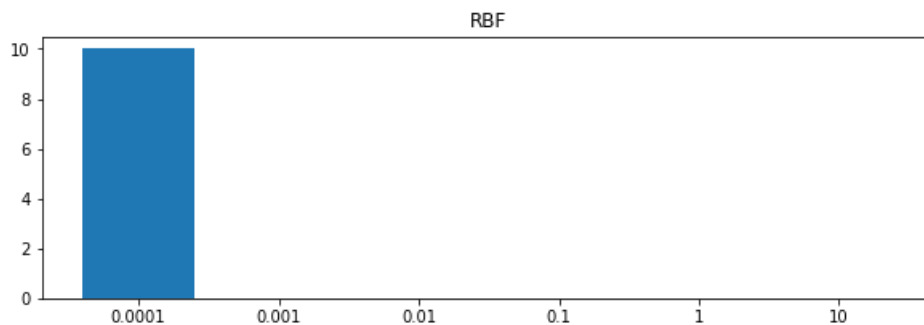
Linear

optimal lambda: [0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001]



RBF

optimal lambda: [0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001]



```

In [10]: optimal_lambda = lamb_list[op_lbin.index(max(op_lbin))]
cvals = []
for i in range(10):
    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)
    cvals.append(cross_validate(x2d, y2d, linear_transform, shuffled_indices, lamb=optimal_lambda))

opt_mean = np.mean(cvals)
opt_std = np.std(cvals)
print("Linear:\n Optimal Lambda:",optimal_lambda,"\n MSE mean:",opt_mean,"\n MSE stdv:", opt_std)

roptimal_lambda = lamb_list[rop_lbin.index(max(rop_lbin))]
rcvals = []
for i in range(10):
    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)
    rcvals.append(cross_validate(x2d, y2d, rbf_transform, shuffled_indices, lamb=roptimal_lambda))
opt_mean = np.mean(rcvals)
opt_std = np.std(rcvals)
print("\nRBF:\n Optimal Lambda:",roptimal_lambda,"\n MSE mean:",opt_mean,"\n MSE stdv:", opt_std)

```

```

Linear:
Optimal Lambda: 0.0001
MSE mean: 0.10938421693716593
MSE stdv: 0.0017477982036840799

RBF:
Optimal Lambda: 0.0001
MSE mean: 16.68747424286159
MSE stdv: 0.15326666906242165

```

Adding noise  $\sim N(0, 100)$

```

In [11]: x2d, y2d_ = load_dataset('dataLinReg2D.txt')
ye = np.random.normal(0, 10, y2d.shape[0])
y2d = y2d_ + ye

```

```

In [12]: lamb_list = [1e-4, 1e-3, 1e-2, 1e-1, 1, 10]

opt_lamb = []
opt_cval = []

ropt_lamb = []
ropt_cval = []

for k in range(10):
    cval_list = []
    rcval_list = []

    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)

    for l in lamb_list:
        # print(">> Lambda",l)
        cval_list.append(cross_validate(x2d, y2d, linear_transform, shuffled_indices, lamb=l))
        rcval_list.append(cross_validate(x2d, y2d, rbf_transform, shuffled_indices, lamb=l))

    opt_idx = cval_list.index(min(cval_list))
    opt_lamb.append(lamb_list[opt_idx])
    opt_cval.append(min(cval_list))

    ropt_idx = rcval_list.index(min(rcval_list))
    ropt_lamb.append(lamb_list[ropt_idx])
    ropt_cval.append(min(rcval_list))

```




```
In [13]: op_lbin = bin_opt_lamb(opt_lamb)
rop_lbin = bin_opt_lamb(ropt_lamb)

print("Linear")
print('optimal lambda:',opt_lamb)
# print('cval losses:', opt_cval)
plot_bin(op_lbin, "Linear")

print("RBF")
print('optimal lambda:',ropt_lamb)
# print('cval losses:', ropt_cval)
plot_bin(rop_lbin, "RBF")
```


Linear

optimal lambda: [0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001]



RBF

optimal lambda: [0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001, 0.0001]



```

In [14]: optimal_lambda = lamb_list[op_lbin.index(max(op_lbin))]
cvals = []
for i in range(10):
    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)
    cvals.append(cross_validate(x2d, y2d, linear_transform, shuffled_indices, lamb=optimal_lambda))

opt_mean = np.mean(cvals)
opt_std = np.std(cvals)
print("Linear:\n Optimal Lambda:",optimal_lambda,"\n MSE mean:",opt_mean,"\n MSE stdv:", opt_std)

roptimal_lambda = lamb_list[rop_lbin.index(max(rop_lbin))]
rcvals = []
for i in range(10):
    shuffled_indices = np.arange(x2d.shape[0])
    np.random.shuffle(shuffled_indices)
    rcvals.append(cross_validate(x2d, y2d, rbf_transform, shuffled_indices, lamb=roptimal_lambda))
opt_mean = np.mean(rcvals)
opt_std = np.std(rcvals)
print("\nRBF:\n Optimal Lambda:",roptimal_lambda,"\n MSE mean:",opt_mean,"\n MSE stdv:", opt_std)

```

Linear:

Optimal Lambda: 0.0001  
MSE mean: 1111.5944621570166  
MSE stdv: 25.546773689208546

RBF:

Optimal Lambda: 0.0001  
MSE mean: 1105.444180170073  
MSE stdv: 12.082771949084563

As can be seen, while the histogram of the optimal lambda doesn't change with the noise, the distribution of the cross-validation error within the optimal lambda adjust to the new values.