

Homework 4

Theory Part

Q1 Consider a layer in CNN that takes in a single channel input of 64×64 , and has 96 filters. In each of the following cases, compute the number of parameters that are learned in this layer. We assume that bias is present for each weight.

[1] A convolution layer with filters of same size as the input.

$$\begin{aligned}Param &= (ksize^2 + 1) * channel \\Param &= (64^2 + 1) * 96 \\Param &= 393312\end{aligned}$$

```
In [1]: (64**2 + 1)* 96
```

```
Out[1]: 393312
```

[2] A convolution layer with 8×8 filters with stride of 4

$$\begin{aligned}Param &= (ksize)^2 * channel + bias \\Param &= (8^2 + 1) * 96 \\Param &= 6240\end{aligned}$$

```
In [2]: (8**2 + 1) * 96
```

```
Out[2]: 6240
```

[3]. A convolution layer with 1×1 filter and a stride of 1

$$\begin{aligned}Param &= (ksize)^2 * channel + bias \\Param &= (1^2 + 1) * 96 \\Param &= 192\end{aligned}$$

Q2 Suppose you would have a neuron which has an RBF kernel as activation function (remember the evil wolf? Drop your linear style of thoughts. Circumferential thoughts can be nice too.)

$$y = \exp(-(x_1^2 + x_2^2)) + b$$

with parameter b . What would be the shapes realized by the set of points $\{(x_1, x_2) : y((x_1, x_2)) = 0\}$ as a function of b ? Explain in at most 2 sentences and/or a little math.

$$0 = \exp(-(x_1^2 + x_2^2)) + b$$

$$-b = \exp(-(x_1^2 + x_2^2))$$

$$-\ln(-b) = x_1^2 + x_2^2$$

Therefore, it is a circle centered around the origin with radius $\sqrt{-\ln(-b)}$. Obviously this is only valid when $-1 < b < 0$.

Supposed now we add weight,

$$y = \exp(-(w_1 x_1^2 + w_2 x_2^2)) + b$$

what shapes can we realize now? Explain in at most 5 sentences and/or a little math. You can make references to publicly available in the internet materials to explain.

$$-\ln(-b) = w_1 x_1^2 + w_2 x_2^2$$

$$-\ln(-b) = \frac{x_1^2}{w_1^{-1}} + \frac{x_2^2}{w_2^{-1}}$$

$$1 = \frac{x_1^2}{-\ln(-b) \cdot w_1^{-1}} + \frac{x_2^2}{-\ln(-b) \cdot w_2^{-1}}$$

$$1 = \left(\frac{x_1}{\sqrt{-\ln(-b) \cdot w_1^{-1}}} \right)^2 + \left(\frac{x_2}{\sqrt{-\ln(-b) \cdot w_2^{-1}}} \right)^2$$

hence it is an ellipse, centered around the origin with radius $\sqrt{-\ln(-b) \cdot w_1^{-1}}$ along the x_1 axis and radius $\sqrt{-\ln(-b) \cdot w_2^{-1}}$ along the x_2 axis.

Q3 Suppose you have five linear neurons n_1, \dots, n_5 , realizing above decision boundaries as shown in Figure 1. That is: for every decision boundary we have outputs are = 0.5 in the zones marked with red plusses, and = 0.2 in the zones marked with the blue minuses.

figure1

As you know, each neuron is realized by:

$$n_i = 0.3H(w_1^{(i)} x_1 + w_2^{(i)} x_2 + b^{(i)}) + 0.2, \quad H(z) \in \{0, 1\}$$

where H is the threshold activation function. You want to predict positive values in a shape marked in green in Figure 1. You want to achieve this prediction by combining these neurons using a threshold neuron H :

$$y = H\left(\sum_i v_i^* n_i + b^*\right)$$

[1] _what do you have to do with the weights of n_5 so that you can move the decision boundary of n_5 so that you can realize the shape in green shown above (in the sense of having positive values inside and negative values outside.)? Give a qualitative description. Note: Give a qualitative description in 3 sentences at most. Note that there is an x - and an y -axis, which helps you to express vectors qualitatively._

The position of the decision boundary of n_5 depends on its weight and biases. Particularly, the ratio between w_1 and w_2 determines the slant of the boundary, while the ratio between the b and w_2 determines its offset from origin. As the desired position is a shift upwards (given that the boundary continues infinitely), we want to decrease b so that the boundaries shift upwards.

[2] _after moving the decision boundary of n_5 appropriately, the green shape looks a bit like an logical AND-combination of the +-zones for every neuron. How to choose the weights v_i^* and the bias b^* in $y = H(\sum_i v_i^* n_i + b^*)$ so that you can realize the green shape (in the sense of having positive values inside and negative values outside that shape)? Note: n_i gives out values either 0.5 or 0.2_

Lets say that function H has a threshold $h = 0$, such that

$$H(z) = 1[z > h] = 1[z > 0]$$

If we were to take green area as $\{+1\}$, for the threshold neuron to fire $+1$, we will need all neurons n_i to fire 0.5,

$$1 = H(\sum_i 0.5v_i^* + b^*)$$

$$1 = 1[(\sum_i 0.5v_i^* + b^*) > 0]$$

$$(\sum_i 0.5v_i^* + b^*) > 0 \quad - - \quad (1)$$

If we were to take non-green area as $\{-1\}$, for the threshold neuron to fire -1 , we will need at least one neuron n_j to fire 0.2,

$$0 = H(\sum_{i=1} v_i^* n_i + v_j^* n_j + b^*)$$

$$0 = H(\sum_{i=1} 0.5v_i^* + 0.2v_j^* + b^*)$$

$$0 = 1[\sum_{i=1} 0.5v_i^* + 0.2v_j^* + b^* > 0]$$

$$\sum_{i=1} 0.5v_i^* + 0.2v_j^* + b^* \leq 0 \quad - - \quad (2)$$

For simplicity, we set $v_i^* = 1$ for all i . Finding b ,

$$0.5 \cdot 5 + b^* > 0 \quad - - \quad (1)$$

$$2.5 + b^* > 0$$

$$b^* > -2.5$$

$$0.5 \cdot 4 + 0.2 + b^* \leq 0 \quad - - \quad (2)$$

$$2.2 + b^* \leq 0$$

$$b^* \leq -2.2$$

$$therefore \quad -2.5 < b \leq -2.2$$

we can pick any b within this range, e.g. $b = -2.4$ with our $v_i^* = 1$, such that

$$\begin{aligned} H(\sum_i 0.5 - 2.4) &= H(2.5 - 2.4) \quad \text{for all } n_i = 0.5 \\ &= H(0.1) = 1[0.1 > 0] = 1 \quad (\text{green}) \end{aligned}$$

$$\begin{aligned} H(\sum_{i=1} 0.5 + 0.2 - 2.4) &= H(2.2 - 2.4) \quad \text{for four } n_i = 0.5 \text{ and one } n_j = 0.2 \\ &= H(-0.2) = 1[-0.2 > 0] = 0 \quad (\text{outside}) \end{aligned}$$

Coding - Part 1

For this section, there is a script called `hw4_code.py` that will hold some definitions e.g. the `FlowerDataset` class and `train_model` function.

Below we will do some sanity check on the dataset.

```

In [3]: from hw4_code import *
import matplotlib.pyplot as plt
%matplotlib inline
from torchvision import transforms

# testing dataset
flower_dataset = FlowerDataset('..\datasets\flowersstuff\102flowers\flowers_data', mode='train')
flower_dataset_val = FlowerDataset('..\datasets\flowersstuff\102flowers\flowers_data', mode='val')

for i in range(3):
    flower1 = flower_dataset[i]
    print('label:', flower1['label'])
    image = transforms.ToPILImage()(flower1['image'])

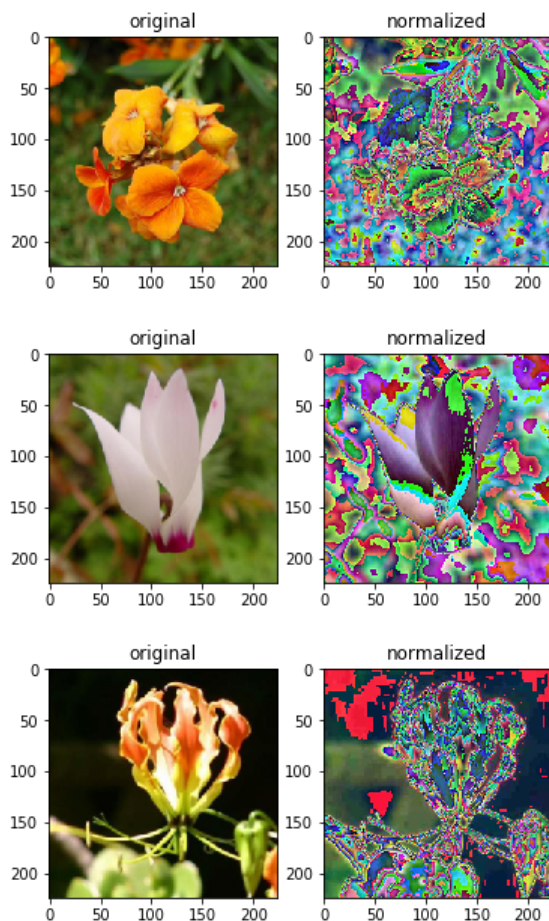
    invTrans = transforms.Compose([ transforms.Normalize(mean = [ 0., 0., 0. ],
                                                    std = [ 1/0.229, 1/0.224, 1/0.225 ]),
                                transforms.Normalize(mean = [ -0.485, -0.456, -0.406 ],
                                                    std = [ 1., 1., 1. ]),
                                transforms.ToPILImage(),
                                ])

    clear_image = invTrans(flower1['image'])

    plt.figure()
    plt.subplot(121)
    plt.imshow(clear_image)
    plt.title('original')
    plt.subplot(122)
    plt.imshow(image)
    plt.title('normalized')

```

label: 45
label: 87
label: 20



We will then use resnet models to classify our flowers.

```
In [19]: # Getting pretrained resnet
from torchvision import models

def get_pretrained_resnet(use_gpu=True):
    model = models.resnet18(pretrained=True)
    if use_gpu:
        model = model.cuda(0)
    return model

resnetmodel = get_pretrained_resnet()
resnet_dict = resnetmodel.state_dict()
```

```
In [20]: # training with various Learn rate
for lr in [0.1, 1, 10]:
    print("\n\nWith learning rate {},\n".format(lr))
    learnrate = lr
    optimizer = optim.SGD(resnetmodel.parameters(), lr=learnrate, momentum=0.9)

    trainedmodel = train_model(flower_dataset, resnetmodel, optimizer,
                               epoch=5, mode='train', use_gpu=True, print_every=1)
    trainedmodel = train_model(flower_dataset_val, trainedmodel, optimizer,
                               epoch=1, mode='val', use_gpu=True, print_every=1)
```

With learning rate 0.1,

```
[train] - Epoch 0..
>> Epoch loss 1.84200 accuracy 0.032 in 436.1029s
[train] - Epoch 1..
>> Epoch loss 1.19028 accuracy 0.197 in 447.7514s
[train] - Epoch 2..
>> Epoch loss 0.84594 accuracy 0.354 in 437.3213s
[train] - Epoch 3..
>> Epoch loss 0.63037 accuracy 0.500 in 438.2014s
[train] - Epoch 4..
>> Epoch loss 0.48444 accuracy 0.625 in 437.8016s
[val] - Epoch 0..
>> Epoch loss 0.50749 accuracy 0.564 in 30.7290s
```

With learning rate 1,

```
[train] - Epoch 0..
>> Epoch loss 0.32693 accuracy 0.737 in 438.2095s
[train] - Epoch 1..
>> Epoch loss 0.06966 accuracy 0.976 in 438.3346s
[train] - Epoch 2..
>> Epoch loss 0.02628 accuracy 0.999 in 437.9225s
[train] - Epoch 3..
>> Epoch loss 0.01591 accuracy 1.000 in 436.8100s
[train] - Epoch 4..
>> Epoch loss 0.01158 accuracy 1.000 in 437.1951s
[val] - Epoch 0..
>> Epoch loss 0.11473 accuracy 0.881 in 29.1180s
```

With learning rate 10,

```
[train] - Epoch 0..
>> Epoch loss 0.12158 accuracy 0.903 in 438.5465s
[train] - Epoch 1..
>> Epoch loss 0.32397 accuracy 0.666 in 440.6043s
[train] - Epoch 2..
>> Epoch loss 0.05289 accuracy 0.945 in 437.9545s
[train] - Epoch 3..
>> Epoch loss 0.00945 accuracy 0.993 in 437.7154s
[train] - Epoch 4..
>> Epoch loss 0.00301 accuracy 0.999 in 437.9883s
[val] - Epoch 0..
>> Epoch loss 0.04652 accuracy 0.950 in 30.8740s
```

In [21]: # now lets do the same thing, but train with an empty resnet

```
def get_empty_resnet(use_gpu=True):
    model = models.resnet18(pretrained=False)
    if use_gpu:
        model = model.cuda(0)
    return model

emptymodel = get_empty_resnet()
empty_dict = resnetmodel.state_dict()
```

In [22]: # training with various Learn rate

```
for lr in [0.1, 1, 10]:
    print("\n\nWith learning rate {},\n".format(lr))
    learnrate = lr
    optimizer = optim.SGD(emptymodel.parameters(), lr=learnrate, momentum=0.9)

    trainedmodel = train_model(flower_dataset, emptymodel, optimizer,
                               epoch=5, mode='train', use_gpu=True, print_every=1)
    trainedmodel = train_model(flower_dataset_val, trainedmodel, optimizer,
                               epoch=1, mode='val', use_gpu=True, print_every=1)
```

With learning rate 0.1,

```
[train] - Epoch 0..
>> Epoch loss 1.64123 accuracy 0.019 in 437.0512s
[train] - Epoch 1..
>> Epoch loss 1.44798 accuracy 0.039 in 436.9365s
[train] - Epoch 2..
>> Epoch loss 1.33094 accuracy 0.050 in 437.3202s
[train] - Epoch 3..
>> Epoch loss 1.25442 accuracy 0.065 in 436.8402s
[train] - Epoch 4..
>> Epoch loss 1.19877 accuracy 0.071 in 436.8860s
[val] - Epoch 0..
>> Epoch loss 1.15961 accuracy 0.076 in 28.8649s
```

With learning rate 1,

```
[train] - Epoch 0..
>> Epoch loss 1.08050 accuracy 0.087 in 436.6231s
[train] - Epoch 1..
>> Epoch loss 0.97139 accuracy 0.114 in 436.3192s
[train] - Epoch 2..
>> Epoch loss 0.91377 accuracy 0.144 in 435.9403s
[train] - Epoch 3..
>> Epoch loss 0.86773 accuracy 0.166 in 437.0819s
[train] - Epoch 4..
>> Epoch loss 0.82263 accuracy 0.205 in 436.9051s
[val] - Epoch 0..
>> Epoch loss 0.84350 accuracy 0.240 in 29.9855s
```

With learning rate 10,

```
[train] - Epoch 0..
>> Epoch loss 0.97167 accuracy 0.123 in 438.8269s
[train] - Epoch 1..
>> Epoch loss 0.79522 accuracy 0.222 in 440.8582s
[train] - Epoch 2..
>> Epoch loss 0.66263 accuracy 0.326 in 439.2992s
[train] - Epoch 3..
>> Epoch loss 0.54665 accuracy 0.427 in 437.4545s
[train] - Epoch 4..
>> Epoch loss 0.44037 accuracy 0.529 in 438.5057s
[val] - Epoch 0..
>> Epoch loss 0.74751 accuracy 0.394 in 30.1826s
```

```
In [29]: # now lets do the same thing, but train with an unfrozen resnet
import torch.nn as nn

def get_unfrozen_resnet(numcl, use_gpu=True):
    model = models.resnet18(pretrained=False)
    for param in model.parameters():
        param.requires_grad = False

    num_fts = model.fc.in_features
    model.fc = nn.Linear(num_fts, numcl)

    if use_gpu:
        model = model.cuda(0)
    return model

unfrozenmodel = get_unfrozen_resnet(102)
unfrozen_dict = resnetmodel.state_dict()
```

```
In [31]: # training with various learn rate
for lr in [0.1, 1, 10]:
    print("\n\nWith learning rate {},\n".format(lr))
    learnrate = lr
    # optimizing only for fully-connected layers
    optimizer = optim.SGD(unfrozenmodel.fc.parameters(), lr=learnrate, momentum=0.9)

    trainedmodel = train_model(flower_dataset, unfrozenmodel, optimizer,
                               epoch=5, mode='train', use_gpu=True, print_every=1)
    trainedmodel = train_model(flower_dataset_val, trainedmodel, optimizer,
                               epoch=1, mode='val', use_gpu=True, print_every=1)
```

With learning rate 0.1,

```
[train] - Epoch 0..
>> Epoch loss 1.15228 accuracy 0.028 in 121.3677s
[train] - Epoch 1..
>> Epoch loss 1.12945 accuracy 0.037 in 120.7015s
[train] - Epoch 2..
>> Epoch loss 1.12416 accuracy 0.045 in 122.1937s
[train] - Epoch 3..
>> Epoch loss 1.12162 accuracy 0.047 in 121.4762s
[train] - Epoch 4..
>> Epoch loss 1.11976 accuracy 0.051 in 121.7667s
[val] - Epoch 0..
>> Epoch loss 1.11822 accuracy 0.047 in 27.9434s
```

With learning rate 1,

```
[train] - Epoch 0..
>> Epoch loss 1.12177 accuracy 0.052 in 121.6749s
[train] - Epoch 1..
>> Epoch loss 1.11096 accuracy 0.057 in 121.2603s
[train] - Epoch 2..
>> Epoch loss 1.10197 accuracy 0.060 in 121.2079s
[train] - Epoch 3..
>> Epoch loss 1.09436 accuracy 0.062 in 121.2120s
[train] - Epoch 4..
>> Epoch loss 1.08763 accuracy 0.066 in 122.1703s
[val] - Epoch 0..
>> Epoch loss 1.07620 accuracy 0.086 in 28.6345s
```

With learning rate 10,

```
[train] - Epoch 0..
>> Epoch loss 1.14607 accuracy 0.049 in 121.1715s
[train] - Epoch 1..
>> Epoch loss 1.10701 accuracy 0.065 in 121.6716s
[train] - Epoch 2..
>> Epoch loss 1.07449 accuracy 0.078 in 121.7347s
[train] - Epoch 3..
>> Epoch loss 1.04668 accuracy 0.090 in 122.4045s
[train] - Epoch 4..
>> Epoch loss 1.02227 accuracy 0.103 in 122.5093s
[val] - Epoch 0..
>> Epoch loss 1.03797 accuracy 0.128 in 27.9617s
```


Coding Part 2

```
In [6]: from hw4_code2 import *
import numpy as np
```




```
In [7]: print("from samplestr.txt:")
data_tr, label_tr = from_text('samplestr.txt')
plot_dataset(data_tr, label_tr)
```

```
from samplestr.txt:
samples with positive class: 2000
samples with negative class: 8000
```



```
In [9]: print("from sampleste.txt:")
data_te, label_te = from_text('sampleste.txt')
plot_dataset(data_te, label_te)
```

```
from sampleste.txt:
samples with positive class: 2000
samples with negative class: 8000
```



```
In [10]: '''
Creating a new Neural Network
'''
from torch import nn
import torch.nn.functional as F

class OurNet(nn.Module):
    def __init__(self):
        super(OurNet, self).__init__()
        self.fc = nn.Linear(2, 1)

    def forward(self, x):
        x = F.relu(self.fc(x))
        return x

our_net = OurNet()
print(our_net)

OurNet(
  (fc): Linear(in_features=2, out_features=1, bias=True)
)
```

In [11]: *# testing the Dataset subclass from hw4_code2*

```
train_set = ImbaDataset('.', mode='train')
print('train:', train_set[2])
test_set = ImbaDataset('.', mode='test')
print('test:', test_set[2])
```

```
train: {'data': tensor([ 0.7959,  0.3896], dtype=torch.float64), 'label': tensor([ 0.], dtype=torch.float64)}
test: {'data': tensor([-0.5324, -0.7683], dtype=torch.float64), 'label': tensor([ 0.], dtype=torch.float64)}
```

```

In [12]: # simple training
import torch
from torch.utils.data import DataLoader
import torch.optim as optim
from torch.autograd import Variable
from torch.utils.data.sampler import WeightedRandomSampler

our_net = OurNet().double()
optimizer = optim.SGD(our_net.parameters(), lr=0.1)
criterion = nn.BCEWithLogitsLoss()

def train_ournet(dataset, model, optimizer, criterion, batch_sampler=None,
                 mode='train', balanced_acc=True):
    print_truecount = True

    if batch_sampler:
        loader = DataLoader(dataset, batch_sampler=batch_sampler)
    else:
        loader = DataLoader(dataset, batch_size=128, shuffle=True)

    running_loss = 0
    running_corrects = 0
    total_data = 0
    # for balanced sampling
    true_pos = 0
    true_neg = 0
    total_positive = 0
    total_negative = 0

    if mode == 'train':
        model.train(True)
    else:
        model.eval()

    for inputdata in loader:
        if mode == 'train': optimizer.zero_grad()

        data = Variable(inputdata['data'])
        labels = Variable(inputdata['label']).view(-1,1)
        if print_truecount:
            print('First batch true count:', labels.sum().item(), 'out of', len(labels),
                  "({}%)".format(labels.sum().item() * 100 / len(labels)))
            print_truecount = False

        outputs = model.forward(data)
        predictions = (outputs > 0)

        # balanced stuff
        true_pos += (predictions.double() * labels.double()).sum().item()
        true_neg += ((1-predictions).double() * (1-labels).double()).sum().item()
        total_positive += labels.sum().item()
        total_negative += (1-labels).sum().item()

        loss = criterion(outputs, labels)
        if mode == 'train':
            loss.backward()
            optimizer.step()

        # balanced stuff
        # total_data += (labels * (positive_weight-1) + 1).sum().item()
        total_data += len(labels)
        corrects = (predictions.double().cpu() == labels.cpu()).double()
        # weights = (positive_weight - 1) * labels.cpu() + 1

        running_corrects += (corrects).sum().item()
        running_loss += loss.item()

    # print("total data:", total_data)
    running_corrects /= (total_data)#float(128*100)
    if mode == 'train':
        print("Trained with accuracy {} and loss {}".format(running_corrects, running_loss))
    elif mode == 'test':
        print("Tested with accuracy {} and loss {}".format(running_corrects, running_loss))

    if balanced_acc:

```

```

        true_pos_rate = true_pos/total_positive
        true_neg_rate = true_neg/(total_negative)
        b_acc = (true_pos_rate + true_neg_rate) / 2
#         print("    True Positive:", true_pos)
#         print("    True Negative:", true_neg)
        print("    True Positive Rate:", true_pos_rate)
        print("    True Negative Rate:", true_neg_rate)
        print("    Balanced Accuracy:", b_acc)
        print('-----')
    return model

```

In [18]: `our_net = OurNet().double()`

```

our_net_trained = train_ournet(train_set, our_net, optimizer, criterion)
with torch.no_grad():
    our_net_trained = train_ournet(train_set, our_net_trained, optimizer, criterion, mode='test')
    our_net_trained = train_ournet(test_set, our_net_trained, optimizer, criterion, mode='test')
print(our_net_trained.state_dict())

```

```

plot_dataset(data_tr, label_tr)
draw_model(our_net_trained)

```

First batch true count: 23.0 out of 128 (17.96875%)

Trained with accuracy 0.8257 and loss 52.962640983636604

True Positive Rate: 0.5

True Negative Rate: 0.907125

Balanced Accuracy: 0.7035625

First batch true count: 21.0 out of 128 (16.40625%)

Tested with accuracy 0.8257 and loss 52.917739253768154

True Positive Rate: 0.5

True Negative Rate: 0.907125

Balanced Accuracy: 0.7035625

First batch true count: 30.0 out of 128 (23.4375%)

Tested with accuracy 0.8326 and loss 52.92446715943566

True Positive Rate: 0.5


True Negative Rate: 0.91575

Balanced Accuracy: 0.707875

OrderedDict([('fc.weight', tensor([[-0.0862, 0.5305]], dtype=torch.float64)), ('fc.bias', tensor([-0.2884], dtype=torch.float64))])

samples with positive class: 2000

samples with negative class: 8000



To achieve 50-50 on minibatches, we need a Sampler subclass.

In [15]: **from random import shuffle**

```
class OurBatchSampler(object):
    """
    Special batch sampler class that ensures a 50-50
    dataset classes.
    """

    def __init__(self, dataset, batch_size, iteration, random=False, start_shuffle=True):
        self.dataset = dataset
        self.batch_size = batch_size
        self.iteration = iteration
        self.random = random

        self.idx_plus = [idx for idx in range(len(dataset)) if dataset[idx]['label'].item() > 0]
        self.idx_minus = [idx for idx in range(len(dataset)) if idx not in self.idx_plus]

        if start_shuffle:
            shuffle(self.idx_plus)
            shuffle(self.idx_minus)

    def __iter__(self):
        batch = []
        for i in range(self.iteration):
            if self.random:
                idplus = np.random.choice(self.idx_plus, self.batch_size // 2).tolist()
                idminus = np.random.choice(self.idx_minus, self.batch_size // 2).tolist()
            else:
                idplus = [self.get_positive_det() for i in range(self.batch_size // 2)]
                idminus = [self.get_negative_det() for i in range(self.batch_size // 2)]
            batch += idplus + idminus

            yield batch
            batch = []

    def get_positive_det(self):
        idx = self.idx_plus[0]
        self.idx_plus = self.idx_plus[1:] + [idx]
        return idx

    def get_negative_det(self):
        idx = self.idx_minus[0]
        self.idx_minus = self.idx_minus[1:] + [idx]
        return idx

    def __len__(self):
        return self.batch_size * self.iteration

train_sampler = OurBatchSampler(train_set, 128, 100)
print('train sampler built')
test_sampler = OurBatchSampler(test_set, 128, 100)
print('test sampler built')
```

train sampler built
test sampler built

```

In [16]: our_net = OurNet().double()
print('\nWithout 50-50 sampling:')
our_net_trained = train_ournet(train_set, our_net, optimizer, criterion)
# print('Train: accuracy')
# our_net_trained = train_ournet(train_set, our_net_trained, optimizer,
#                               criterion, mode='test')

with torch.no_grad():
    print('Test: accuracy')
    our_net_trained = train_ournet(test_set, our_net_trained, optimizer,
                                   criterion, mode='test')


    plot_dataset(data_tr, label_tr)
    draw_model(our_net_trained)

```

```

Without 50-50 sampling:
First batch true count: 30.0 out of 128 (23.4375%)
Trained with accuracy 0.6439 and loss 53.94670419843951
  True Positive Rate: 0.5
  True Negative Rate: 0.679875
  Balanced Accuracy: 0.5899375
-----
Test: accuracy
First batch true count: 38.0 out of 128 (29.6875%)
Tested with accuracy 0.638 and loss 54.12464868764891
  True Positive Rate: 0.5
  True Negative Rate: 0.6725
  Balanced Accuracy: 0.5862499999999999
-----
samples with positive class: 2000
samples with negative class: 8000

```



```

In [17]: our_net = OurNet().double()
print('\n\n')
print('With 50-50 sampling:')
our_net_trained = train_ournet(train_set, our_net, optimizer, criterion,
                               batch_sampler=train_sampler)

with torch.no_grad():
    print('Test: accuracy')
    our_net_trained = train_ournet(test_set, our_net_trained, optimizer,
                                   criterion, mode='test')
    plot_dataset(data_tr, label_tr)
    draw_model(our_net_trained)
    print(our_net_trained.state_dict())

```

```

With 50-50 sampling:
First batch true count: 64.0 out of 128 (50.0%)
Trained with accuracy 0.775859375 and loss 75.1650439681061
  True Positive Rate: 1.0
  True Negative Rate: 0.55171875
  Balanced Accuracy: 0.775859375

```

```

-----
Test: accuracy
First batch true count: 27.0 out of 128 (21.09375%)
Tested with accuracy 0.6422 and loss 90.2322079772061
  True Positive Rate: 1.0
  True Negative Rate: 0.55275
  Balanced Accuracy: 0.776375

```

```

-----
samples with positive class: 2000
samples with negative class: 8000
OrderedDict([('fc.weight', tensor([[ 0.5646,  0.1262]], dtype=torch.float64)), ('fc.bias', tensor([-0.3454],
dtype=torch.float64))])

```

